

```
import pandas as pd
import numpy as np

path = "/content/drive/MyDrive/framingham.csv"

df = pd.read_csv(path)
df.head()
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHy
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	
3	0	61	3.0	1	30.0	0.0	0	
4	0	46	3.0	1	23.0	0.0	0	

Next steps:

Generate code with df

 View recommended plots

```
df.shape

(4238, 16)

df.size

67808

df.columns

Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',
      'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
      'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'],
      dtype='object')

df.dtypes

male          int64
age           int64
education     float64
currentSmoker int64
cigsPerDay    float64
BPMeds        float64
prevalentStroke int64
prevalentHyp  int64
diabetes      int64
totChol       float64
sysBP         float64
diaBP         float64
BMI           float64
heartRate     float64
glucose       float64
TenYearCHD    int64
dtype: object

df.isnull().sum()

male          0
age           0
education    105
currentSmoker 0
cigsPerDay    29
BPMeds        53
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       50
sysBP         0
diaBP         0
BMI           19
heartRate     1
glucose      388
```

```
TenYearCHD      0
dtype: int64
```

```
df[:2]
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHy
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	

```
df['education'].value_counts(normalize=True)*100
```

```
1.0    41.616259
2.0    30.316961
3.0    16.622308
4.0    11.444471
Name: education, dtype: float64
```

```
from sklearn.impute import KNNImputer
```

```
columns = [['education', 'cigsPerDay', 'BPMeds', 'totChol', 'BMI', 'heartRate', 'glucose']]
impute = KNNImputer(n_neighbors = 9)
for i in columns:
    df[i] = impute.fit_transform(df[i])
```

```
df.isna().sum()
```

```
male          0
age           0
education     0
currentSmoker 0
cigsPerDay    0
BPMeds        0
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       0
sysBP        0
diaBP        0
BMI          0
heartRate    0
glucose      0
TenYearCHD   0
dtype: int64
```

```
#EDA
```

```
df[:2]
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHy
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	

```
df.rename(columns={"male":"gender"}, inplace = True)
```

```
df[:2]
```

	gender	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalent
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	

```
df['gender'].value_counts()
```

```

0    2419
1    1819
Name: gender, dtype: int64

```

```

df['age'].min(), df['age'].max()

(32, 70)

```

```


import matplotlib.pyplot as plt
import seaborn as sns

```

```

sns.histplot(data = df, x="age",kde=True)
plt.show()

```



```

df['currentSmoker'].value_counts(normalize = True)*100


0    50.589901
1    49.410099
Name: currentSmoker, dtype: float64

```

```

sns.histplot(data=df, x="cigsPerDay")
plt.show()

```



```

df['BPMeds'].value_counts()

```

```
0.000000    4100
1.000000    124
0.111111     14
Name: BPMeds, dtype: int64
```

```
y = df['TenYearCHD']
```

```
X = df.drop(columns="TenYearCHD")
```

```
X.shape
```

```
(4238, 15)
```

```
y.shape
```

```
(4238,)
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train.shape
```

```
(3390, 15)
```

```
y_train.shape
```

```
(3390,)
```

```
X_test.shape, y_test.shape
```

```
((848, 15), (848,))
```

```
df[:3]
```

	gender	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalent
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	

```
# prompt: import standard scaler
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
y_pred[:10]
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
y_test[:10]
```

```
3188  0
764    0
3264  0
1967  0
2185  0
393    0
2333  1
```

```
1159    0
3788    0
1674    1
Name: TenYearCHD, dtype: int64
```

```
pd.DataFrame({
    "Actual": y_test,
    "predicted": y_pred
})[:15]
```

	Actual	predicted
3188	0	0
764	0	0
3264	0	0
1967	0	0
2185	0	0
393	0	0
2333	1	0
1159	0	0
3788	0	0
1674	1	0
759	0	0
1803	0	0
410	0	0
157	0	0
3886	0	0

```
model.score(X_test, y_test)
```

```
0.8573113207547169
```

```
model.score(X_train, y_train)
```

```
0.8548672566371681
```

```
# prompt: import r2 score
```

```
from sklearn.metrics import r2_score
```

```
r2_score_model = r2_score(y_test, y_pred)
print("R2 Score:", r2_score_model * 100)
```

```
R2 Score: -14.293352343610755
```

```
model.coef_
```

```
array([[ 0.21494318,  0.57183733,  0.03328953, -0.01271295,  0.27893144,
         0.04401033,  0.07357614,  0.0907214 ,  0.00156521,  0.04087167,
         0.30882841,  0.01788746, -0.01694299, -0.0657164 ,  0.18885262]])
```

```
model.intercept_
```

```
array([-1.97123212])
```

```
# prompt: import accuracy score
```

```
from sklearn.metrics import accuracy_score
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy * 100)
```

```
Accuracy: 85.73113207547169
```

Let's Improve the Score by checking with the CV - Cross validation

```
df1 = df
```

```
df1.head()
```

	gender	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalent
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	
3	0	61	3.0	1	30.0	0.0	0	
4	0	46	3.0	1	23.0	0.0	0	

Next steps: [Generate code with df1](#) [View recommended plots](#)

```
X = df1.drop(columns={"TenYearCHD"})
```

```
y = df['TenYearCHD']
```

```
X.shape, y.shape
```

```
((4238, 15), (4238,))
```

```
# prompt: import train test split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train_cv, X_test, y_train_cv, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
X_train, X_val, y_train, y_val = train_test_split(X_train_cv, y_train_cv, test_size=0.25, random_state=1)
```

```
X_test.shape, y_test.shape
```

```
((848, 15), (848,))
```

[Generate](#) [Using ...](#) [Close](#)

< 1 of 4 > [Undo changes](#) [Use code with caution](#)

```
# prompt: import standard scaler
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
scaler.fit(X_train)
```

```
X_train = scaler.transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
X_val = scaler.transform(X_val)
```

[Generate](#) [Using ...](#) [Close](#)

< 1 of 4 > [Undo changes](#) [Use code with caution](#)

```
# prompt: import Logistic regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(X_train, y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```



```
model.coef_  
  
array([[ 0.22227067,  0.58535055,  0.02438671, -0.06570819,  0.30855561,  
         0.03117776,  0.08007633, -0.00478253,  0.03896179,  0.08912729,  
         0.31058245,  0.08798583, -0.11481075, -0.05661747,  0.22148162]])
```

```
model.intercept_  
  
array([-2.03205876])
```

```
y_pred_cv = model.predict(X_val)  
y_pred_cv[:15]  
  
array([0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
y_val[:15]  
  
1430    1  
1810    0  
1655    0  
2760    0  
4085    0  
4084    0  
1266    1  
3685    0  
2295    0  
2954    1  
1683    0  
1753    0  
3161    0  
2859    0  
4144    1  
Name: TenYearCHD, dtype: int64
```

```
pd.DataFrame({  
    "Actual": y_val,  
    "predicted": y_pred_cv  
}[:35])
```

	Actual	predicted	
1430	1	0	
1810	0	0	
1655	0	0	
2760	0	0	
4085	0	0	
4084	0	1	
1266	1	0	
3685	0	0	
2295	0	0	
2954	1	0	
1683	0	0	
1753	0	0	
3161	0	0	
2859	0	0	
4144	1	0	
1746	0	0	
2613	0	0	
3220	1	0	
1815	0	0	
2059	0	0	
636	0	0	
0	0	0	
2974	0	0	
1105	0	0	
1544	0	0	
3094	1	0	
2443	1	0	
323	0	0	
3073	0	0	
1625	0	0	
2810	0	0	

Checking how much varies from the Data


```

zero = 0
one = 0
for i in y_pred_cv:
    if i == 0:
        zero +=1
    else:
        one +=1

print(zero)
print(one)

827
21

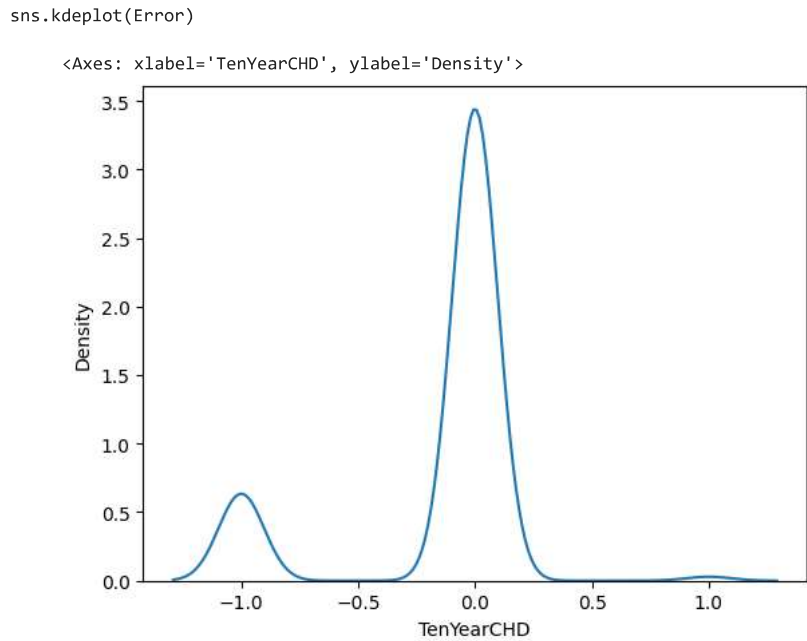
# Error
y_val.shape, y_pred_cv.shape

((848,), (848,))

Error = y_pred_cv - y_val
Error[:10]

1430  -1
1810   0
1655   0
2760   0
4085   0
4084   1
1266  -1
3685   0
2295   0
2954  -1
Name: TenYearCHD, dtype: int64

```



Check for the MSE

Generate

Using ...

import mean sqaure error

Q

Close

<

1 of 3

>

Undo changes

Use code with caution

prompt: import mean sqaure error

```

from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_val, y_pred_cv)
print("Mean squared error:", mse)

Mean squared error: 0.1615566037735849

```